

VaxNeRF

Roger Mari
roger.mari@ens-paris-saclay.fr

MLBriefs presentation day, 12 december 2023

Reminder: NeRF

- What is a **Neural Radiance Field (NeRF)** ?
 - Continuous function f that encodes a static 3D scene. [Mildenhall et al. 2020]
 - f is learned by one or more MLPs F_{θ} .

$$f(x, d) \rightarrow (c, \sigma)$$

Reminder: NeRF

$$f(x, d) \rightarrow (c, \sigma)$$

- What is a **Neural Radiance Field (NeRF)** ?

- Continuous function f that encodes a static 3D scene. [Mildenhall et al. 2020]

- f is learned by one or more MLPs F_{θ} .

- Input: $x \rightarrow$ spatial coords of a 3D point + $d \rightarrow$ view-specific features (e.g. viewing dir, embedding vecs)

- Output: $c \rightarrow$ RGB color + $\sigma \rightarrow$ geometry magnitude (e.g. volume density, sdf)

- *Training data:*



N views

+ N camera models (typically pinhole)

Reminder: NeRF

- The NeRF training strategy:


- Render the color $C(\mathbf{r})$ of a ray $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$ with near and far bounds t_n and t_f :

$$C(\mathbf{r}) = \sum T_i \alpha_i c_i$$

α = opacity = degree in which light can pass through

T = transparency: degree of visibility of the point

$$f(\mathbf{x}, \mathbf{d}) \rightarrow (\mathbf{c}, \sigma)$$



**“IMPRESSIVE QUALITY
AND DETAILS”**



Cahiers du Cinéma
(ne pas confondre
avec “Chair du CVPR”)



OFFICIAL SELECTION
FESTIVAL DE CANNES

**“TOO SLOW,
I FELL ASLEEP”**



Reviewer 2

**“YOU WILL LOVE ITS
UNSUPERVISED NATURE”**



Reviewer 1

**“UNREALISTIC.
TOO MANY POINTS OF VIEW.”**



Reviewer 3



1. Convert a NeRF MLP to a different data structure

- Only improves inference time.
- E.g. sparse voxel octrees (PlenOctrees, Neural Sparse Voxel Fields).

2. Subdivide the scene (2021)

- Each region of space is assigned an individual MLP
- E.g. Uniform or hierarchical subdivisions (KiloNeRF, Recursive-NeRF)
- Learnable subdivisions (DeRF, non-regular Voronoi decomposition)

3. Voxel grids that cache density + learnable features (2022)

- DVGO or Plenoxels optimize neural features or spherical harmonics for each voxel.
- Shallow MLP or none at all.

4. Multi-resolution hash grids (2022)

- Instant-NGP. Also for aerial imagery in GridNeRF or Neuralangelo.

5. Gaussian splatting (2023)

Other tools

JaxNeRF


NerfAcc

Nerfstudio



The VaxNeRF Method (Voxel-Accelerated NeRF)

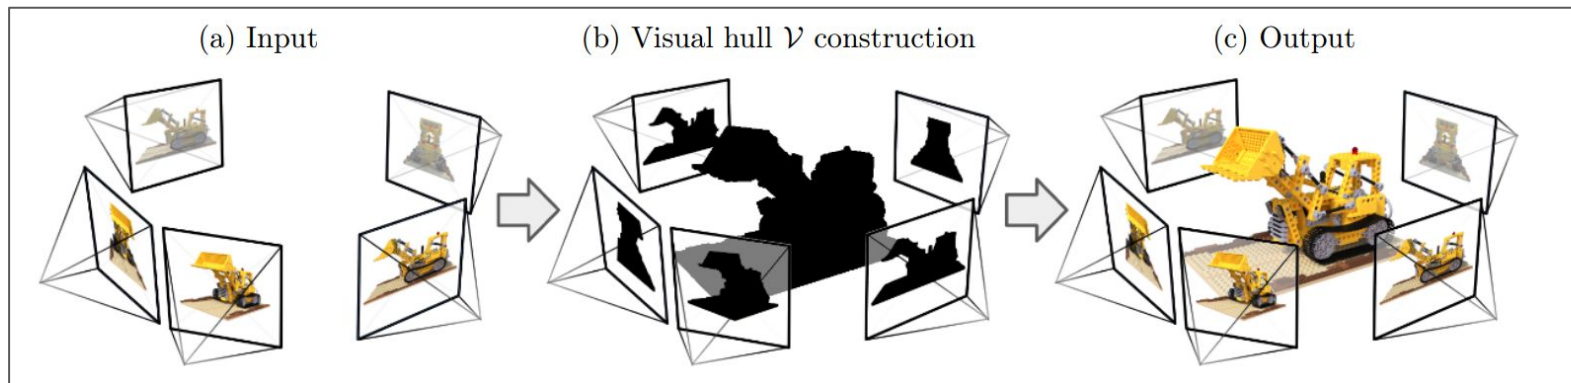
The NeRF search space can be reduced by exploiting the **visual hull**.




The visual hull is the intersection of the set of foreground silhouettes back-projected into the 3D space. Toy example using 3 input cameras.

The VaxNeRF Method (Voxel-Accelerated NeRF)


VaxNeRF uses this insight to dramatically reduce the optimization time. [Kondo et al. 2021]



NeRF
(2020)



VaxNeRF
(2021)



Kondo, Naruya, et al. "Vaxnerf: Revisiting the classic for voxel-accelerated neural radiance field." *arXiv preprint* (2021).

The VaxNeRF Method (Voxel-Accelerated NeRF)

VaxNeRF Advantages

- *Integration is very simple. Requires minimal code modifications.*
- *Training is 2-8x faster with respect to the original NeRF.*
- *No loss of quality in the output.*
- *Outperforms coarse-fine sampling strategies using 2 MLPs.*

VaxNeRF Disadvantages

- *Foreground-background segmentation may not be feasible or straight-forward.*
- *A 3D array to represent a visual hull requires more memory than an MLP.
(e.g. 400x400x400 voxels ~70MB w.r.t. ~7MB for an MLP)*
- *Gain of speed depends on the visual hull size.*



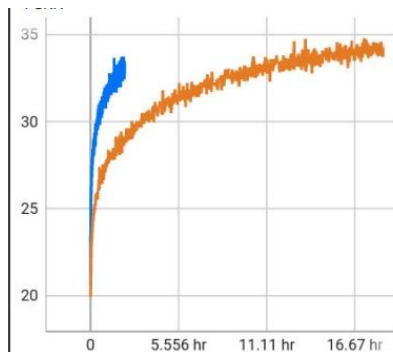

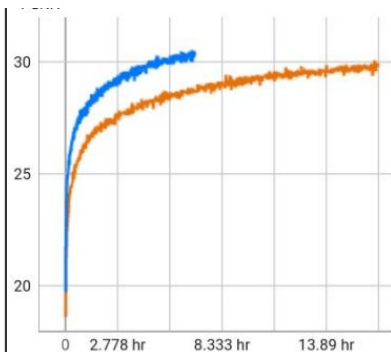
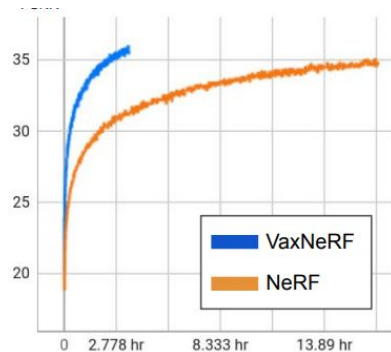
Misleading name issue: the name voxel-accelerated NeRF (VaxNeRF) may not be the most descriptive. The visual hull could be represented in a format other than a voxel grid.

Results

Classic NeRF: 2 MLP
(64 and 128 pts per ray)

VaxNeRF: 1 MLP
(600 pts per ray)

PSNR vs optimization time



Visual hull
(~10 s to
compute)




Results

“Lego” scene: 100 training views of 800x800px

Visual hull:



< 1 min



VaxNeRF:



5 min

Results

“Ficus” scene: 100 training views of 800x800px

Visual hull:



< 1 min

VaxNeRF:



5 min

Results



Visual Hull



0.5min



1min



NeRF




VaxNeRF

5min

Generate your own custom novel views from pretrained models

Select input(s) 📘 Upload data



Parameters Réinitialiser

| | | | |
|---------------|---|------|--|
| scene | <div>lego</div> | | Scene |
| azimuth | <div><div></div></div> | 59,5 | Max: 360 Min: 0 Azimuth angle in degrees |
| elevation | <div><div></div></div> | 32,6 | Max: 90 Min: -90 Elevation angle in degrees |
| radius | <div><div></div></div> | 1 | Max: 1.5 Min: 0.5 Distance to the scene |
| scaling | <div><div></div></div> | 0,5 | Max: 1 Min: 0.25 Scaling factor (e.g. 0.5 = downsample x2) |
| model | <input checked="" type="radio"/> VaxNeRF <input type="radio"/> NeRF | | Model architecture |
| training step | <input type="radio"/> 2000 <input checked="" type="radio"/> 512000 | | Checkpoint step to load |

Run

